General hints:

- Your code should work with *Python 3.10*.
- We suggest to adhere to the PEP8 style guide and use line lengths of up to 100. One way to achieve this is to format your code with black: `black {source_file_or_dir} -l 100` or add black to your IDE.
- `for` loops can be slow in Python, use vectorized `numpy` operations wherever possible (see assignment 1 for an example).

How to run the exercise and tests:

- See the setup instructions downloadable from our website for installation details.
- We always assume you run commands in the *root folder* of the exercise.
- If you're using miniconda, don't forget to activate your environment with `conda activate cvenv`
- Install the required packages with `pip install -U -r requirements.txt`
- Python files in the *root folder* of the repository contain the scripts to run the code.
- Some exercises use jupyter notebooks. It should be already installed from the requirements. Run `jupyter lab` to start the server.
- Some exercises contain unittests in folder `tests/` to check your solution. Run `python -m pytest .` to run all tests.

In this course, you will learn about Computer Vision. In this first exercise you will set do some small exercises to brush up your *linear algebra* and to get familiar with *python* and *numpy*.

At this point you should have worked through the `setup.pdf` and have a working python 3.8 installation ready. You should know how to navigate and run commands in the command prompt.

1. Pen and Paper tasks

Now you'll brush up your linear algebra a little by doing eigendecomposition by hand. Later we will see how we can do the same in python.

If you'd like to review Linear Algebra, we recommend Khan Academy. The Matrix Cookbook is another useful resource.

1) Calculate the eigendecomposition of the following matrix:

$$A = \begin{bmatrix} 9 & 1 \\ 8 & 7 \end{bmatrix}$$

2) Use the eigendecomposition of A above to show how you can efficiently compute A^{10} (you don't have to show the final value of the matrix).

3) You are given the following matrix:

$$A = \begin{bmatrix} 4 & 8 & 2 \\ 8 & 41 & 24 \\ 2 & 24 & 21 \end{bmatrix}$$

$\det(A) = 400$
One of the eigen values is 1. Find the other two. **Hint:** You don't have to calculate the eigenvalues from scratch. Use the properties of eigenvalues.

4) You are given the following matrix:

$$A = \begin{bmatrix} 100 & 100 & 100 \\ 99 & 99 & 102 \\ 98 & 98 & 104 \end{bmatrix}$$

$\det(A) = 0$
Find the eigenvalues of A.

2. Code Warmup

To solve code exercises, you have to fill in code between the *START TODO* and *END TODO* markers in the code. Before you run any code, make sure you have the correct conda environment activated, and don't forget to install the required packages with `pip install -U -r requirements.txt` .

Todo: In the file `lib/example_file.py`, complete `example_function` by adding

```
return input_variable * 2
```

Execute the command `python example_script.py` to see the function in action.

Run `python -m tests.test_example` to test if the function is implemented correctly.

3. Getting to know numpy

1) Numpy tensors

You will now play around with some basics of tensor manipulation in *numpy*. The basic object in numpy is an homogeneous multidimensional array. Numpy's array class is called *ndarray*. Here is a quickstart tutorial: https://numpy.org/devdocs/user/quickstart.html

While working with numpy it will be crucial to understand indexing of tensors, since this is a very basic concept that will reappear in most exercises: https://numpy.org/doc/stable/user/basics.indexing.html

Todo: Run the script `run_numpy_arrays.py`. We will walk you through it's code and output during this exercise.

Let's create two matrices and check their properties.

```
A = np.array(np.arange(4))
B = np.array([-1, 3])
print(f"A (shape: {A.shape}, type: {type(A)}) = {A}")
print(f"B (shape: {B.shape}, type: {type(B)}) = {B}")
```

Output:

```
A (shape: (4,), type: <class 'numpy.ndarray'>) = [0 1 2 3]
B (shape: (2,), type: <class 'numpy.ndarray'>) = [-1  3]
```

First, 2 arrays (also called tensors in the context of deep learning) are created. Each numpy tensor has an attribute `numpy.ndarray.shape` which describes the dimensions of the defined tensor. Type, shape and content of the tensors are the first output of the script. Please note how we are using f-strings to output variables.

In order to perform matrix multiplication and addition in numpy there are two methods: numpy.matmul and numpy.add. Please read their respective documentation in numpy before proceeding.

Next, we try to multiply the two tensors with `matmul`.

```
np.matmul(A, B)
```

Output:

```
ValueError: matmul: Input operand 1 has a mismatch in its core dimension 0,
with gufunc signature (n?,k),(k,m?)->(n?,m?) (size 2 is different from 4)
```

We get a *ValueError* due to the shape mismatch between the two numpy arrays we want to multiply/add. In order to deal with different array shapes during arithmetic operations, we can either reshape the arrays or broadcast the smaller array across the larger one such that they have compatible shapes.

```
C = A.reshape([2, 2])
```

```
print(f"C shape: {C.shape}, content:\n{C}")
```

Output:

```
C shape: (2, 2), content:
[[0 1]
 [2 3]]
```

Now the matrix multiplication CB works out.

```
matmul_result = np.matmul(C, B)
print(matmul_result)
```

Output:

```
[3 7]
```

When adding the C with shape $(2, 2)$ and B with shape $(2,)$, B will be automatically broadcast to match the shape of C.

```
print(np.add(C, B))
```

Output:

```
[[-1  4]
 [ 1  6]]
```

The star operator $*$ will do an element-wise multiplcation between the C and B. Again, B will be broadcasted to fit.

```
print(C * B)
```

Output:

```
[[ 0  3]
 [-2  9]]
```

The function np.diag can transform the vector B shaped $(2,)$ into a diagonal matrix of shape $(2, 2)$.

```
print(np.diag(B))
```

Output:

```
[[-1  0]
 [ 0  3]]
```

For transposing a ndarray use numpy.transpose or the method `numpy.ndarray.T`.

```
print(np.transpose(C))
```

Output:

```
[[0 2]
 [1 3]]
```

Tensor operations are a central part of the exercises and deep learning in general, so play around with the script to get familiar with them. You could also just start python with the command `python` and play around in there.

2) **Remember that `for` loops can be slow**

Use vectorized numpy expressions instead of manual loops wherever possible. The following is an example for computing the sum $\sum_{i=0}^{N-1} i^2$ with $N = 1000000$:

This is **wrong** (Takes about *200ms*).

```
total = 0
for i in range(1000000):
    total += i ** 2
```

This is **correct** (Takes about *8ms*):

```
import numpy as np
numbers = np.arange(1000000, dtype=np.int64)
total = (numbers ** 2).sum()
```

Note: The `np.int64` datatype means we use integers that are large enough to store the result. The square and sum operations are vectorized and run in fast C code internally.

3) **Eigendecomposition**

Using numpy.linalg you can also perform many linear algebra functionalities. Given a square and symmetric matrix A, the eigendecomposition $A = Q\Lambda Q^T$ with $\Lambda = diag(\lambda_1, \ldots, \lambda_n)$ can be done using numpy.linalg.eig.

Todo: Run the script `run_eigen.py` to see the eigendecomposition in action for $A = \begin{bmatrix} 7 & -\sqrt{3} \\ -\sqrt{3} & 5 \end{bmatrix}$

Do not worry about the `NotImplementedError`, you will fix that now.

Todo: In file `lib/eigendecomp.py`, complete the function `get_matrix_from_eigdec` to return the square matrix A, given its eigenvalues $\lambda_1, \ldots, \lambda_n$ and eigenvectors Q as an input. **Note:** We are using type hints to make the code more readable and give you hints about the input and output.

Todo: Complete the `get_euclidean_norm` and `get_dot_product` functions. These take vectors as input and are used to show that the columns of Q are orthonormal, i.e. the columns are of unit length and the columns are orthogonal (their dot product is 0).

Todo: Complete the `get_inverse` function by using that $A^{-1} = Q\Lambda^{-1}Q^T$ with $\Lambda^{-1} = diag(\lambda_1^{-1}, \ldots, \lambda_n^{-1})$ is the inverse of A. Do **not** use `numpy.linalg.inv`. You can invert the diagonal matrix Λ without it.

4) **Vector Norms**

The length of a vector is not a single number but can be defined in different ways. These vector norms share common properties but also have different characteristics. In numpy you can use the numpy.linalg.norm function to compute the L_p norm of a numpy array, where $p \geq 1$.

More formally, the L_p norm of a vector x is defined as: $\|x\|_p = \left(\sum_{i=1}^{n} |x_i|^p\right)^{1/p}$

For $p = 1$ we get the Manhattan norm, for $p = 2$ we get the Euclidian norm and for $p = \infty$ we approximate the maximum norm: $\|x\|_\infty = \max_i |x_i|$.

Now to plot the norms, we create a 2-dimensional grid and color the norm values as a heat map. To create the grid, we use meshgrid which can transform two 1-dimensional vectors into a grid representation. For example:

```
N = 3
lin_x = np.linspace(-1, 1, N)  # e.g. for N=3: [-1, 0, 1] with shape (3,)
lin_y = np.linspace(-1, 1, N)  # same
X, Y = np.meshgrid(lin_x, lin_y)
print(f"X (shape {X.shape}):\n{X}\n")
print(f"Y (shape {Y.shape}):\n{Y}\n")
```

Output:

```
X (shape (3, 3)):
[[-1.  0.  1.]
 [-1.  0.  1.]
 [-1.  0.  1.]]

Y (shape (3, 3)):
[[-1. -1. -1.]
 [ 0.  0.  0.]
 [ 1.  1.  1.]]
```

This way, you can get the 2D-coordinates of the grid at i, j by accessing $X_{i,j}$ and $Y_{i,j}$

Todo: in file `lib/norms.py`, complete the function `get_norm` to compute the norm of each 2D vector composed by the i, j-th elements of the matrices X and Y. You will first need to stack the two matrices together to form a (N, N, 2) tensor (using np.stack) and then compute the L_p-norm (using np.linalg.norm). This way, you get the norm results for each point in the grid. Run the file `plot_norms.py` to see a plot of the 2-norm.

We have added an argument to the script with the argparse package. Run the command `python plot_norms.py -p 1` to see a plot of the 1-norm.

4. Diffusion filters

Run the jupyter notebook `diffusion.ipynb` and complete the missing parts. See the top of this file on how to run jupyter notebooks.

This assignment will be introduced on 21.10. 16:00 CET. The solution will be released and discussed in the following session.

Thanks to the team of the course Foundations of Deep Learning, Machine Learning Lab, University of Freiburg for providing parts 1 to 3 of this exercise.